

CST8502 - Lab 5

Python Exercise

Student Name: Alvin Litani Liauw

Student Number: 041118874

For every step, include screenshot of the code and the results in this document (screenshot from colab/jupyter notebook). Also, in your words, explain your code and results. If there is no explanation, no marks will be given. No need to write long paragraphs, but one or 2 lines per step.

1

Read CSV file and put it inside dataframe

```
df = pd.read_csv('train.csv')
```

Show number of attributes

```
print(f'Number of attributes: {df.shape[1]}')
```

Show names of attributes

```
print(f'Name of attributes: {df.columns}')
```

Show number of instances

```
print(f'Number of instances: {df.shape[0]}')
```

Show columns and top 5 rows

```
df.head()
```

Number of attributes: 12

Name of attributes: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'], dtype='object')

Number of instances: 891

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

2

```
# Check for duplicates after dropping PassengerId
df = df.drop(columns='PassengerId')
print(f'Number of duplicated rows: {df.duplicated().sum()}')

# Show only columns with missing data
missing_data = df.isnull().sum()
print(f'{missing_data[missing_data > 0]}')

# Check columns that have more than 50% missing data
missing_proportion = df.isnull().mean()
columns_to_drop = missing_proportion[missing_proportion > 0.5]
```

```
Number of attributes: 7
Name of attributes: Index(['Survived', 'Pclass', 'Sex', 'Fare', 'Embarked', 'AgeGroup',
                          'Relatives'],
                          dtype='object')
```

3

```
# Drop those columns together with "Name" and "Ticket" columns
df = df.drop(columns=['Name', 'Ticket'] + list(columns_to_drop.index))
```

```
Number of duplicated rows: 0
Age          177
Cabin        687
Embarked      2
dtype: int64
```

4

```
# Bin Age column into AgeGroup
df['AgeGroup'] = df['Age'].apply(lambda x:
                                'NK' if pd.isnull(x) else
                                'Child' if x < 16 else
                                'Youth' if x < 30 else
                                'Adult' if x < 65 else
                                'Senior')
```

```
# Define the bins and labels
```

```
bins = [-float('inf'), 0, 3, float('inf')]
labels = ['None', 'Few', 'Many']
```

```
# Create new 'Relatives' column based on number of relatives
df['Relatives'] = pd.cut(df['Parch'] + df['SibSp'], bins=bins, labels=labels)
```

```
df['Fare'] = df['Fare'].apply(lambda x:
                               'Free' if x == 0 else
                               'Low' if x < 50 else
                               'Average' if x < 100 else
                               'High')
```

```
# Drop SibSp, Parch, Age columns
```

```
df = df.drop(columns=['SibSp','Parch','Age'])
print(f'Number of attributes: {df.shape[1]}')
print(f'Name of attributes: {df.columns}')
```

```
Number of attributes: 7
Name of attributes: Index(['Survived', 'Pclass', 'Sex', 'Fare', 'Embarked', 'AgeGroup',
                          'Relatives'],
                          dtype='object')
```

5

```
# Do one-hot encoding
```

```
data_encoded = pd.get_dummies(df,drop_first = True)
data_encoded.head()
```

	Survived	Pclass	Sex_male	Fare_Free	Fare_High	Fare_Low	Embarked_Q	Embarked_S	AgeGroup_Child	AgeGroup_NK	AgeGroup_Senior	AgeGroup_Youth	Relatives
0	0	3	True	False	False	True	False	True	False	False	False	True	
1	1	1	False	False	False	False	False	False	False	False	False	False	
2	1	3	False	False	False	True	False	True	False	False	False	True	
3	1	1	False	False	False	False	False	True	False	False	False	False	
4	0	3	True	False	False	True	False	True	False	False	False	False	

6

```
# Split data into training and testing data
```

```
label = data_encoded['Survived']
attributes = data_encoded.drop(columns='Survived')
```

```
attributes_train, attributes_test, label_train, label_test = train_test_split(attributes, label,
test_size=0.3)
```

```
attributes_train.info()
```

```
label_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 623 entries, 846 to 848
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pclass                623 non-null   int64
1   Sex_male              623 non-null   bool
2   Fare_Free             623 non-null   bool
3   Fare_High             623 non-null   bool
4   Fare_Low              623 non-null   bool
5   Embarked_Q            623 non-null   bool
6   Embarked_S            623 non-null   bool
7   AgeGroup_Child        623 non-null   bool
8   AgeGroup_NK           623 non-null   bool
9   AgeGroup_Senior       623 non-null   bool
10  AgeGroup_Youth         623 non-null   bool
11  Relatives_Few          623 non-null   bool
12  Relatives_Many         623 non-null   bool
dtypes: bool(12), int64(1)
memory usage: 17.0 KB
<class 'pandas.core.series.Series'>
Index: 623 entries, 846 to 848
Series name: Survived
Non-Null Count  Dtype
-----
623 non-null    int64
dtypes: int64(1)
memory usage: 9.7 KB
```

7

```
# Train the model on the training data
```

```
model = DecisionTreeClassifier(max_depth=3)
```

```
model.fit(attributes_train, label_train)
```

8

```
label_pred = model.predict(attributes_test)
```

```
print(f'Number of survivors in testing set: {label_pred.sum()} out of {label_pred.size}')
```

Number of survivors in testing set: 99 out of 268

9

Show accuracy and confusion matrix

```
accuracy = accuracy_score(label_test, label_pred)
```

```
print(f'Accuracy: {accuracy * 100:.2f}%')
```

```
ConfusionMatrix = confusion_matrix(label_test, label_pred)
```

```
print(ConfusionMatrix)
```

Create the tree diagram

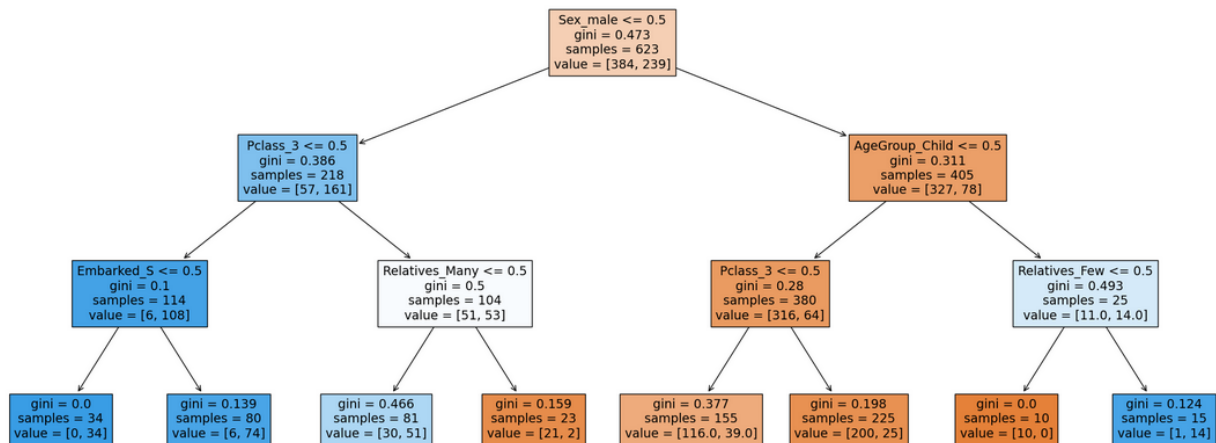
```
fig = plt.figure(figsize=(25,10))
```

```
tree.plot_tree(model, feature_names=attributes_train.columns, filled=True)
```

```
plt.show()
```

Accuracy: 82.46%

```
[[144  21]
 [ 26  77]]
```



10

Split data into training and testing data

```
label = data_encoded['Survived']
```

```
attributes = data_encoded.drop(columns='Survived')
```

```
attributes_train, attributes_test, label_train, label_test = train_test_split(attributes, label,
test_size=0.3)
```

11

```
# Scale the training data
scaler = StandardScaler()
attributes_train = scaler.fit_transform(attributes_train)
attributes_test = scaler.transform(attributes_test)
```

12

```
# Do one-hot encoding
data_encoded = pd.get_dummies(df, drop_first = True)
print(data_encoded)
```

13

```
# Initialize the k-NN classifier with a chosen k value (e.g., k=5)
knn = KNeighborsClassifier(n_neighbors=5)
```

```
# Train the k-NN model on the training data
knn.fit(attributes_train, label_train)
```

```
# Predict on the test set
label_pred = knn.predict(attributes_test)
```

```
# Show accuracy and confusion matrix for testing set
accuracy = accuracy_score(label_test, label_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')
```

```
ConfusionMatrix = confusion_matrix(label_test, label_pred)
print(ConfusionMatrix)
```

```
Accuracy: 76.87%
[[135  33]
 [ 29  71]]
```

14

```
# Standardize the data
scaler = StandardScaler()
df_scaled = scaler.fit_transform(data_encoded)
```

```


# Sum of squared distances for each k
inertia = []

# Range of k values to test
k_values = range(1, 16)

# Perform K-Means for each k and store the inertia
for k in k_values:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(df_scaled)
    inertia.append(kmeans.inertia_)

# Plot the elbow graph
plt.figure(figsize=(10, 5))
plt.plot(k_values, inertia, marker='x')
plt.xlabel('Number of Clusters/K')
plt.ylabel('Inertia (Sum of Squared Distances)')
plt.title('Elbow Method for K')
plt.show()
print('Depending on the random seed, k is usually either 7 or 8.')

```



Depending on the random seed, k is usually either 7 or 8.

```

# Perform Local Outlier Factor/LOF with 10% outlier assumption
lof = LocalOutlierFactor(contamination=0.1)
lof_outliers = lof.fit_predict(df_scaled)
lof_outliers = np.where(lof_outliers == -1)[0] # LOF labels outliers as -1

# Perform Isolation Forest/ISF with 10% outlier assumption
isf = IsolationForest(contamination=0.1)
isf_outliers = isf.fit_predict(df_scaled)
isf_outliers = np.where(isf_outliers == -1)[0] # ISF labels outliers as -1

# Find common outliers
common_outliers = np.intersect1d(lof_outliers, isf_outliers)

# Print common outliers
print("Number of common outliers:", len(common_outliers))
print("Common outliers:\n", df.iloc[common_outliers])

```

Number of common outliers: 16

Common outliers:

	Survived	Pclass	Sex	Age	Fare	Embarked	Relatives
16	0	3	male	2.0	29.1250	Q	5
116	0	3	male	70.5	7.7500	Q	0
122	0	2	male	32.5	30.0708	C	1
171	0	3	male	4.0	29.1250	Q	5
188	0	3	male	40.0	15.5000	Q	2
245	0	1	male	44.0	90.0000	Q	2
278	0	3	male	7.0	29.1250	Q	5
280	0	3	male	65.0	7.7500	Q	0
301	1	3	male	28.0	23.2500	Q	2
303	1	2	female	28.0	12.3500	Q	0
322	1	2	female	30.0	12.3500	Q	0
361	0	2	male	29.0	27.7208	C	1
412	1	1	female	33.0	90.0000	Q	1
626	0	2	male	57.0	12.3500	Q	0
787	0	3	male	8.0	29.1250	Q	5
885	0	3	female	39.0	29.1250	Q	5